Portfolio Assignment 1, Methods 3, 2021, autumn semester

Study Group 8, Emma Rose Hahn (EH), Viara Krasteva (VK), Kristian Nøhr Villebro (KV), Luke Ring (LR)

29/09/2021

Assignment 1: Using mixed effects modelling to model hierarchical data

In this assignment we will be investigating the *politeness* dataset of Winter and Grawunder (2012) and apply basic methods of multilevel modelling.

```
#this is the dataset we will be exploring in this exercise
politeness <- read.csv('politeness.csv') ## read in data
```

Exercise 1 - describing the dataset and making some initial plots

1. Describe the dataset, such that someone who happened upon this dataset could understand the variables and what they contain

Exercise 1, part 1– (EH)

The politeness dataset contains the data obtained from the study of Korean formal and informal speech (https://doi.org/10.1016/j.wocn.2012.08.006) which investigated the fundamental frequency of male and female participants' speech in a variety of formal and informal scenarios.

The following table describes the variables in the dataset:

Variable	Description
`subject`	participant ID
`gender`	participant's gender
`scenario`	the experimental scenario from 1 to 7 such as "asking a favour"
`attitude`	either 'inf' for informal stimuli or 'pol' for formal stimuli
`total_duration`	duration of participant's response in seconds
`f0mn`	mean fundamental frequency (f0) of the participant's speech
`hiss_count`	number of times the participants made a noisy breath intake

Remark: The `gender` , `scenario` and `attitude` variables should be encoded as factors as they show a categorical function withing this dataset . In addition, these variables have non-unique values across participants, and are not ordered.

```
#Encoding some of the variables as factors (gender, attitude, and scenario)

politeness$attitude <- as.factor(politeness$attitude)
politeness$gender <- as.factor(politeness$gender)
politeness$scenario <- as.factor(politeness$scenario)
```

Exercise 1, part 2–(EH)

2. Create a new data frame that just contains the subject *F1* and run two linear models; one that expresses *f0mn* as dependent on *scenario* as an integer; and one that expresses *f0mn* as dependent on *scenario* encoded as a factor
    i. Include the model matrices, X from the General Linear Model, for these two models in your report and describe the different interpretations of *scenario* that these entail
    ii. Which coding of *scenario*, as a factor or not, is more fitting?

```
# Create a subset dataframe for subject F1 only
pf1 <- politeness[politeness$subject == "F1", ]
pf1
```

```
##    subject gender scenario attitude total_duration  f0mn hiss_count
## 1       F1      F        1      pol         18.392 214.6          2
## 2       F1      F        1      inf         13.551 210.9          0
## 3       F1      F        2      pol          5.217 284.7          0
## 4       F1      F        2      inf          4.247 265.6          0
## 5       F1      F        3      pol          6.791 210.6          0
## 6       F1      F        3      inf          4.126 285.6          0
## 7       F1      F        4      pol          6.244 251.5          1
## 8       F1      F        4      inf          3.245 281.5          0
## 9       F1      F        5      pol          5.625 229.6          1
## 10      F1      F        5      inf          3.950 250.5          0
## 11      F1      F        6      pol         28.508 181.1          1
## 12      F1      F        6      inf         55.159 229.3          0
## 13      F1      F        7      inf         60.309 219.8          2
## 14      F1      F        7      pol         40.825 175.8          0
```

```
# make model predicting f0mn by scenario (integer)
m1<- lm(f0mn ~ as.integer(scenario), data = pf1)

# get model matrix
mm1 <- model.matrix(m1)

# make model predicting f0mn by scenario (factor)
m2 <- lm(f0mn ~ as.factor(scenario), data = pf1)

# get model matrix
mm2 <- model.matrix(m2)
```

```
Here is the model using "scenario" encoded as an integer
```

```
summary(m1)
```

```
##
## Call:
## lm(formula = f0mn ~ as.integer(scenario), data = pf1)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -44.836 -36.807   6.686  20.918  46.421
##
## Coefficients:
##                      Estimate Std. Error t value Pr(>|t|)
## (Intercept)           262.621     20.616  12.738 2.48e-08 ***
## as.integer(scenario)   -6.886      4.610  -1.494    0.161
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 34.5 on 12 degrees of freedom
## Multiple R-squared:  0.1568, Adjusted R-squared:  0.0865
## F-statistic: 2.231 on 1 and 12 DF,  p-value: 0.1611
```

```
mm1
```

```
##    (Intercept) as.integer(scenario)
## 1            1                    1
## 2            1                    1
## 3            1                    2
## 4            1                    2
## 5            1                    3
## 6            1                    3
## 7            1                    4
## 8            1                    4
## 9            1                    5
## 10           1                    5
## 11           1                    6
## 12           1                    6
## 13           1                    7
## 14           1                    7
## attr(,"assign")
## [1] 0 1
```

And here is the model using “scenario” encoded as a factor

```
summary(m2)
```

```
##
## Call:
## lm(formula = f0mn ~ as.factor(scenario), data = pf1)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -37.50 -13.86   0.00  13.86  37.50
##
## Coefficients:
##                      Estimate Std. Error t value Pr(>|t|)
## (Intercept)            212.75      20.35  10.453  1.6e-05 ***
## as.factor(scenario)2    62.40      28.78   2.168   0.0668 .
## as.factor(scenario)3    35.35      28.78   1.228   0.2591
## as.factor(scenario)4    53.75      28.78   1.867   0.1041
## as.factor(scenario)5    27.30      28.78   0.948   0.3745
## as.factor(scenario)6    -7.55      28.78  -0.262   0.8006
## as.factor(scenario)7   -14.95      28.78  -0.519   0.6195
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 28.78 on 7 degrees of freedom
## Multiple R-squared:  0.6576, Adjusted R-squared:  0.364
## F-statistic:  2.24 on 6 and 7 DF,  p-value: 0.1576
```

```
mm2
```

```
##    (Intercept) as.factor(scenario)2 as.factor(scenario)3 as.factor(scenario)4
## 1            1                    0                    0                    0
## 2            1                    0                    0                    0
## 3            1                    1                    0                    0
## 4            1                    1                    0                    0
## 5            1                    0                    1                    0
## 6            1                    0                    1                    0
## 7            1                    0                    0                    1
## 8            1                    0                    0                    1
## 9            1                    0                    0                    0
## 10           1                    0                    0                    0
## 11           1                    0                    0                    0
## 12           1                    0                    0                    0
## 13           1                    0                    0                    0
## 14           1                    0                    0                    0
##    as.factor(scenario)5 as.factor(scenario)6 as.factor(scenario)7
## 1                     0                    0                    0
## 2                     0                    0                    0
## 3                     0                    0                    0
## 4                     0                    0                    0
## 5                     0                    0                    0
## 6                     0                    0                    0
## 7                     0                    0                    0
## 8                     0                    0                    0
## 9                     1                    0                    0
## 10                    1                    0                    0
## 11                    0                    1                    0
## 12                    0                    1                    0
## 13                    0                    0                    1
## 14                    0                    0                    1
## attr(,"assign")
## [1] 0 1 1 1 1 1 1
## attr(,"contrasts")
## attr(,"contrasts")$`as.factor(scenario)`
## [1] "contr.treatment"
```

Conclusion: The above output shows the difference in model matrices between scenario encoded as an integer and factor. The integer version treats scenario as a continuous variable, whereas the factorized version creates a regression line per scenario.

For this dataset, scenario should be a factor, since the scenarios are not a continuous variable and depending on the prescribed scenario, the participants may have a different f0 (mean fundamental frequency of speech), and we are interested in following the trajectory of f0 across scenarios not as a variable that consistently decreases or increases, but a separate regression line showing the changes in f0 between the 7 different scenarios. And in order, to be able to see that crucial difference we need to consider the 'scenario' variable as a factor when we run a model predicting f0 across scenarios.

Exercise 1, part 3 – (EH)

3. Make a plot that includes a subplot for each subject that has *scenario* on the x-axis and *f0mn* on the y-axis and where points are colour coded according to *attitude*
    i. Describe the differences between subjects

```
politeness %>% ggplot(aes(scenario, f0mn, color = attitude)) +
    geom_point() +
    facet_wrap(vars(subject))
```

```
## Warning: Removed 12 rows containing missing values (geom_point).
```

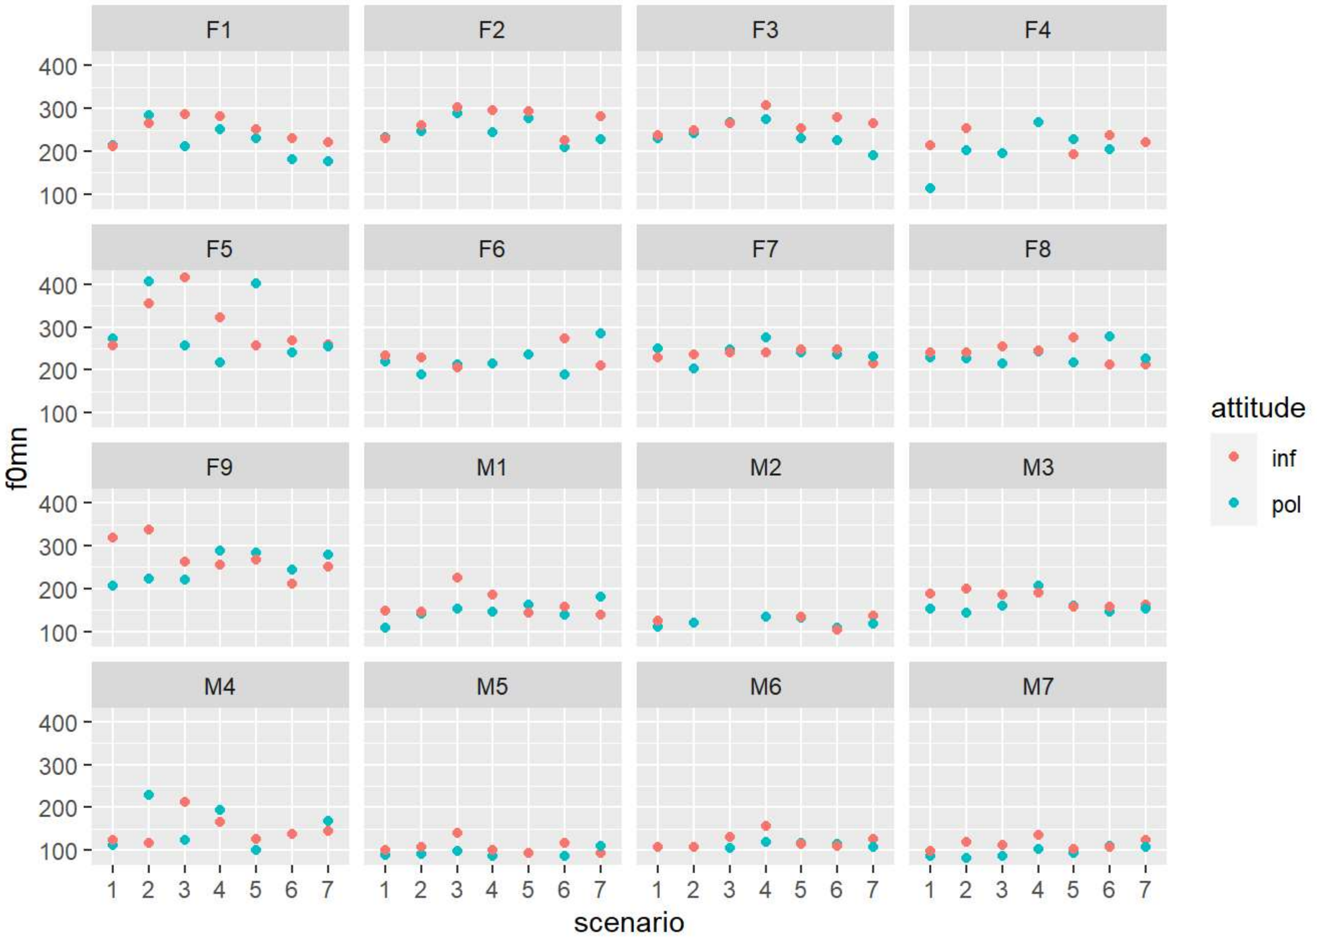

…

We can visually observe that there are baseline differences between the male and female subjects' mean fundamental frequency of speech, where the males' f0 is consistently lower, across scenario and attitude. Between the different scenarios there is variability in the f0 values for both male and female subjects depending on both the scenario type and the attitude (informal or formal). There is a consistent tendency across scenario type and gender for the mean fundamental frequency of speech to be slightly higher when the attitude is informal as opposed to formal. This visual information inerpreted from this plot is consistent with the results of Winter and Grawunder (2012)

##Exercise 2 - comparison of models

```
mixed.model <- lmer(formula=..., data=...)
example.formula <- formula(dep.variable ~ first.level.variable + (1 | second.level.variabl
e))
```

###Part 1 – (VK) 1) Build four different models and do some comparisons

```r
# the single level model
m3 <- lm(formula = f0mn ~ gender, data = politeness)

# a two-level model where each scenario has a unique intercept
m4 <- lmer(formula = f0mn ~ gender + (1 | scenario), data = politeness)
# a two-level model that has models subject as intercept
m5 <- lmer(formula = f0mn ~ gender + (1 | subject), data = politeness)

# a two-level model that incorporate intercepts for both subject and scenario
m6 <- lmer(formula = f0mn ~ gender +
            (1 | subject) + (1 | scenario), data = politeness)

#comparing AIC and Deviance values for all the models
AIC(m3)
```

```
## [1] 2163.971
```

```r
AIC(m4)
```

```
## [1] 2152.314
```

```r
AIC(m5)
```

```
## [1] 2099.626
```

```r
AIC(m6)
```

```
## [1] 2092.482
```

```r
deviance(m3)
```

```
## [1] 327033.6
```

```r
deviance(m4)
```

```
## Warning in deviance.merMod(m4): deviance() is deprecated for REML fits;
## use REMLcrit for the REML criterion or deviance(.,REML=FALSE) for deviance
## calculated at the REML fit
```

```
## [1] 2144.314
```

```r
deviance(m5)
```

```
## Warning in deviance.merMod(m5): deviance() is deprecated for REML fits;
## use REMLcrit for the REML criterion or deviance(.,REML=FALSE) for deviance
## calculated at the REML fit
```

```
## [1] 2091.626
```

```
deviance(m6)
```

```
## Warning in deviance.merMod(m6): deviance() is deprecated for REML fits;
## use REMLcrit for the REML criterion or deviance(.,REML=FALSE) for deviance
## calculated at the REML fit
```

```
## [1] 2082.482
```

```
anova(m4, m5, m6)
```

```
## refitting model(s) with ML (instead of REML)
```

```
## Data: politeness
## Models:
## m4: f0mn ~ gender + (1 | scenario)
## m5: f0mn ~ gender + (1 | subject)
## m6: f0mn ~ gender + (1 | subject) + (1 | scenario)
##    npar    AIC    BIC  logLik deviance   Chisq Df Pr(>Chisq)
## m4    4 2162.3 2175.7 -1077.1   2154.3
## m5    4 2112.1 2125.5 -1052.0   2104.1 50.2095  0
## m6    5 2105.2 2122.0 -1047.6   2095.2  8.8725  1   0.002895 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
piecewiseSEM::rsquared(c(m4, m5, m6))
```

```
##   Response   family     link method  Marginal Conditional
## 1     f0mn gaussian identity   none 0.6779555   0.6967788
## 2     f0mn gaussian identity   none 0.6681651   0.7899229
## 3     f0mn gaussian identity   none 0.6677206   0.8077964
```

The single level model performs the worst and this makes sense as we do not expect all participants to have the same f0 as their voices have naturally occurring differences (not just ones predicted by gender). There are differences that might be explained by the scenario or individual subject (as we observed in the plot above), however, neither of those are taken into account when using a single level model.

So then using two-level models that explain variance by either taking scenario or subject as random intercepts, is definitely an improvement to help explain more of the scenario/attitude based differences, not just the gender differences.

Consequently, Of the three multi-level models,it is model m6, which includes random intercepts for both subject and scenario,that has the most explained variance with for the entire model $R^2 \approx 0.81$ or 81%.

Additionally, we see that model m6 has the lowest AIC and deviance.

Exercise 2, part 2 and 3 – (LR)

2. Why is our single-level model bad?
   i. create a new data frame that has three variables, *subject*, *gender* and *f0mn*, where *f0mn* is the average of all responses of each subject, i.e. averaging across *attitude* and_scenario_
   ii. build a single-level model that models *f0mn* as dependent on *gender* using this new dataset

iii. make Quantile-Quantile plots, comparing theoretical quantiles to the sample quantiles) using `qqnorm` and `qqline` for the new single-level model and compare it to the old single-level model (from 1).i). Which model's residuals (ϵ) fulfil the assumptions of the General Linear Model better?)

iv. Also make a quantile-quantile plot for the residuals of the multilevel model with two intercepts. Does it look alright?

3. Plotting the two-intercepts model

    i. Create a plot for each subject, (similar to part 3 in Exercise 1), this time also indicating the fitted value for each of the subjects for each for the scenarios (hint use `fixef` to get the "grand effects" for each gender and `ranef` to get the subject- and scenario-specific effects)

```
# scenario x f0mn y, attitude = color
ff <- fixef(m6)
ff
```

```
## (Intercept)     genderM
##    246.7650   -115.1746
```

```
rf <- ranef(m6)
rf <- as.data.frame(rf)
rf
```

```
##      grpvar        term grp    condval   condsd
## 1   subject (Intercept)  F1 -10.490356 8.280794
## 2   subject (Intercept)  F2  10.251809 8.280794
## 3   subject (Intercept)  F3   3.795129 8.280794
## 4   subject (Intercept)  F4 -29.495270 9.095403
## 5   subject (Intercept)  F5  47.093999 8.280794
## 6   subject (Intercept)  F6 -18.396273 8.794359
## 7   subject (Intercept)  F7  -6.976691 8.280794
## 8   subject (Intercept)  F8  -8.521934 8.280794
## 9   subject (Intercept)  F9  12.739587 8.280794
## 10  subject (Intercept)  M1  21.052117 8.280794
## 11  subject (Intercept)  M2  -5.462358 9.453009
## 12  subject (Intercept)  M3  33.561535 8.280794
## 13  subject (Intercept)  M4  16.093337 8.524543
## 14  subject (Intercept)  M5 -28.267430 8.280794
## 15  subject (Intercept)  M6 -12.640202 8.794541
## 16  subject (Intercept)  M7 -24.336998 8.280794
## 17 scenario (Intercept)   1 -11.595496 5.488728
## 18 scenario (Intercept)   2   5.321218 5.532205
## 19 scenario (Intercept)   3   6.795658 5.586194
## 20 scenario (Intercept)   4  11.348815 5.578013
## 21 scenario (Intercept)   5   1.411037 5.488705
## 22 scenario (Intercept)   6  -8.622136 5.489258
## 23 scenario (Intercept)   7  -4.659096 5.488058
```

```
politeness$effect_gender <- 0.0
politeness[politeness$gender == "F", ]$effect_gender <- ff[1]
politeness[politeness$gender == "M", ]$effect_gender <- ff[1] + ff[2]
politeness$intercept_subject <- left_join(politeness, rf, by = c("subject" = "grp"), copy
 = TRUE, keep = FALSE)$condval
politeness$intercept_scenario <- left_join(politeness, rf, by = c("scenario" = "grp"), cop
y = TRUE, keep = FALSE)$condval
politeness$predicted <- politeness$effect_gender + politeness$intercept_subject + politene
ss$intercept_scenario
politeness %>% ggplot(aes(scenario, f0mn, color = attitude)) +
    geom_point() +
    geom_point(aes(y = predicted, shape = "fitted values"), color = "black", size = 2) +
    scale_shape_manual(name = "model", values = c(18)) +
    facet_wrap(vars(subject))
```

```
## Warning: Removed 12 rows containing missing values (geom_point).
```

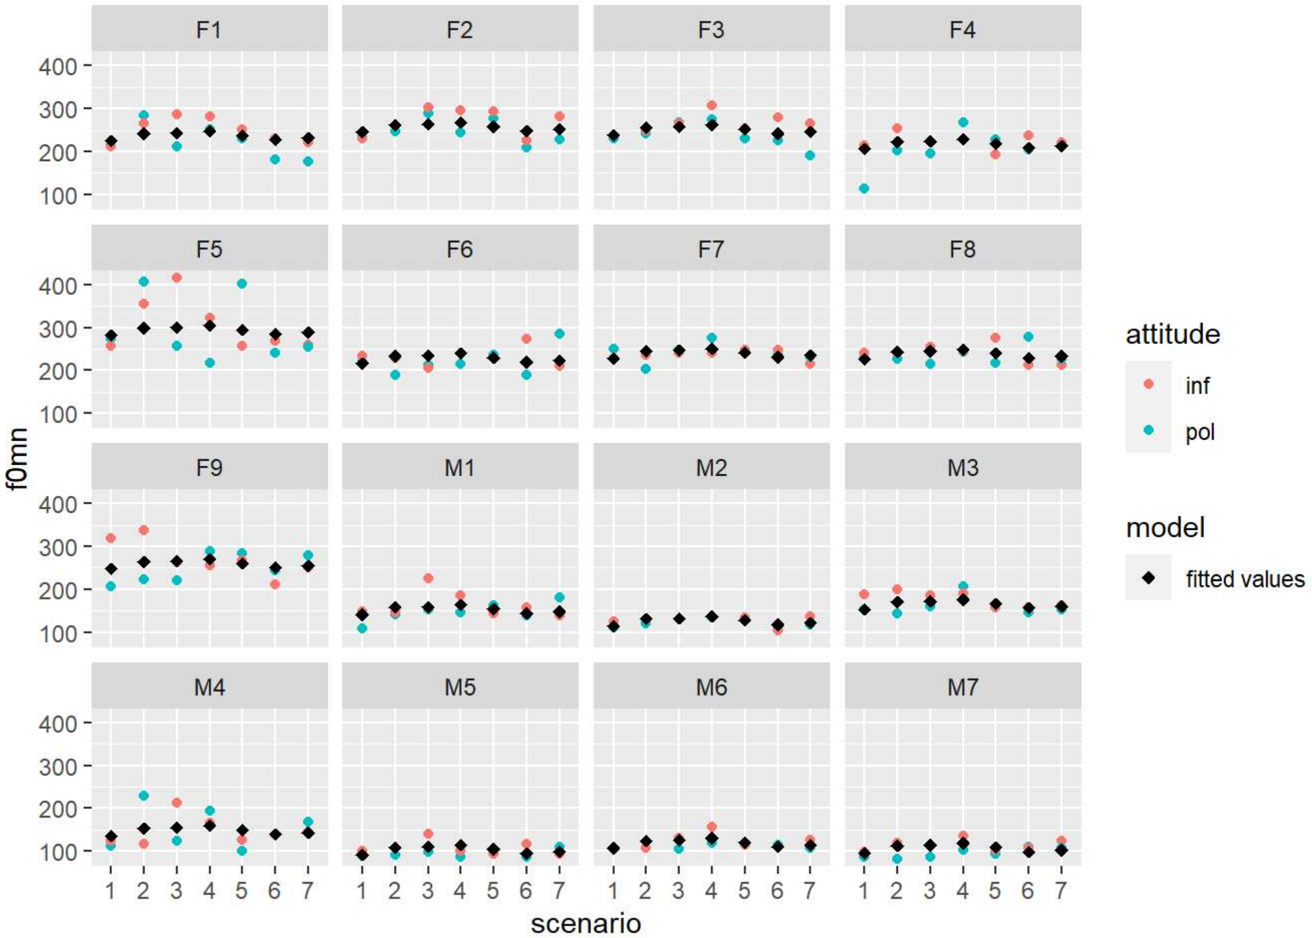


```
deviance(m3)
```

```
## [1] 327033.6
```

```
deviance(m4, REML = FALSE)
```

```
## [1] 2154.33
```

```
deviance(m5, REML = FALSE)
```

```
## [1] 2104.175
```

```
deviance(m6, REML = FALSE)
```

```
## [1] 2095.279
```

```
politeness_aggregated <- politeness[!is.na(politeness$f0mn), ] %>% group_by(subject) %>% s
ummarize(subject = subject[1], gender = gender[1], f0mn = mean(f0mn))
politeness_aggregated
```

```
## # A tibble: 16 x 3
##    subject gender  f0mn
##    <chr>   <fct>  <dbl>
##  1 F1      F       235.
##  2 F2      F       258.
##  3 F3      F       251.
##  4 F4      F       212.
##  5 F5      F       299.
##  6 F6      F       225.
##  7 F7      F       239.
##  8 F8      F       237.
##  9 F9      F       261.
## 10 M1      M       155.
## 11 M2      M       122.
## 12 M3      M       169.
## 13 M4      M       150.
## 14 M5      M       100.
## 15 M6      M       118.
## 16 M7      M       104.
```

```
m7 <- lm(f0mn ~ gender, data = politeness_aggregated)
par(mfrow=c(1,2))
qqnorm(fitted.values(m7))
qqline(fitted.values(m7))
qqnorm(fitted.values(m2))
qqline(fitted.values(m2))
```

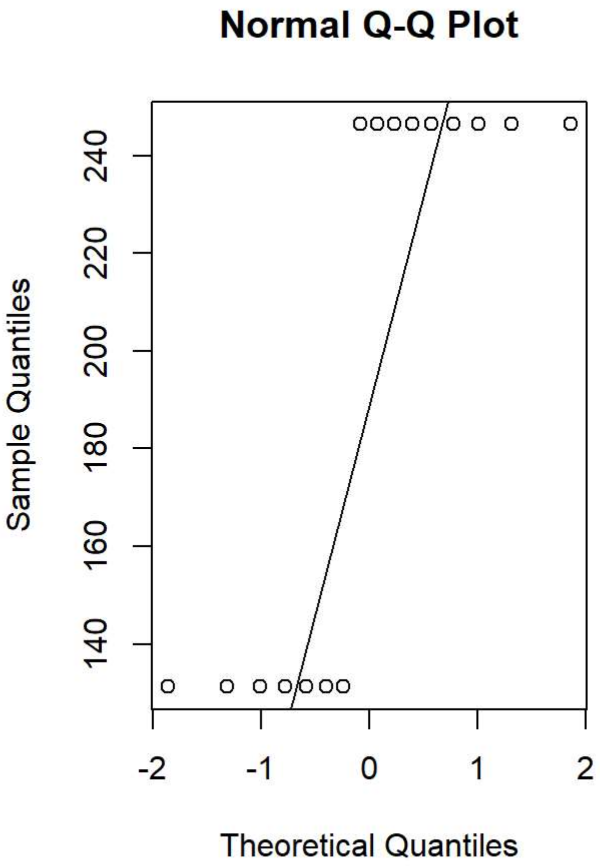

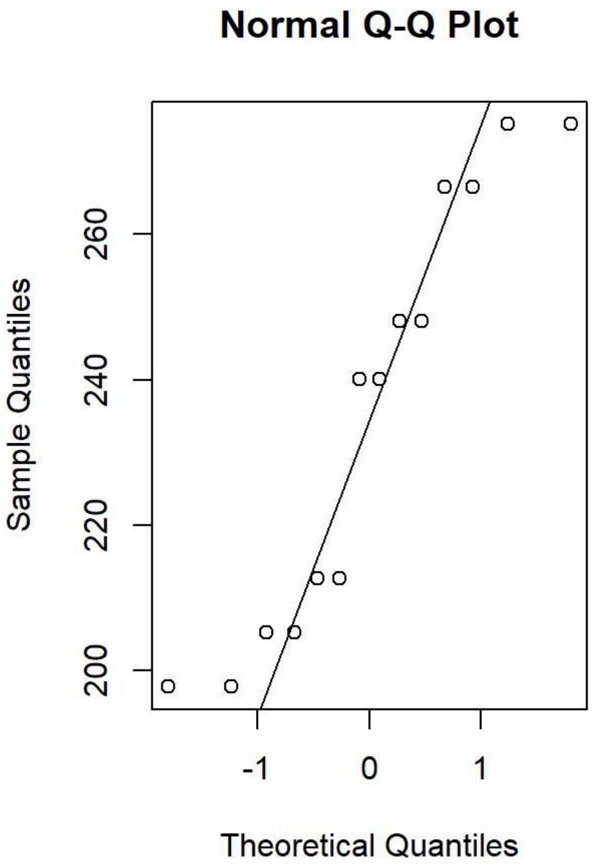


```
par(mfrow=c(1,1))
qqnorm(fitted.values(m6))
qqline(fitted.values(m6), col = "red")
```

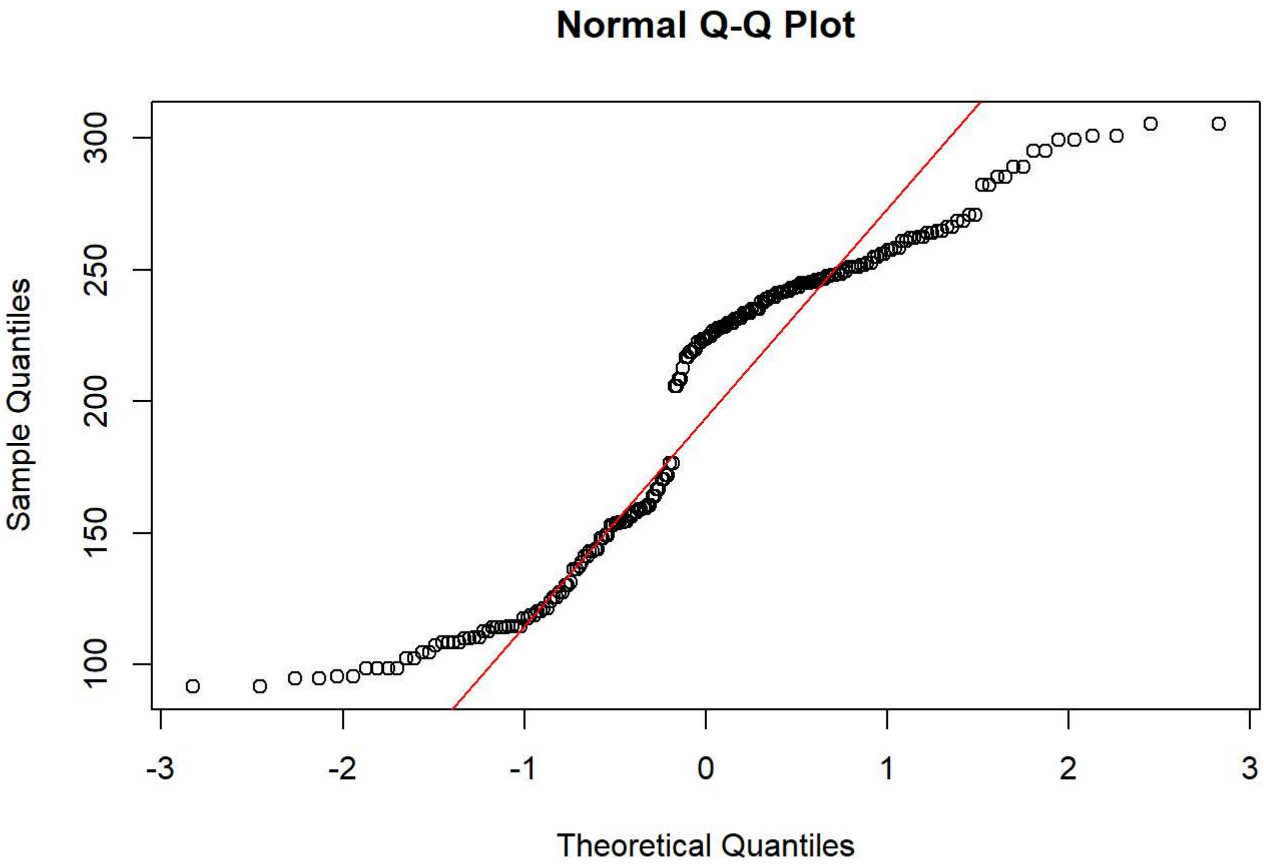

Assessing the QQ-plots of the single-level models it seems that the aggregated model m7's residuals are worse off than those of model m2. The residuals of model m2 are better dispersed along the line - however it still doesn't look fantastic. The QQ-plot of the multilevel model m6 looks better than any of the single level ones, with data points closer to the line and more evenly dispersed on both sides of the line.

Looking at the plot for the observed and the fitted values it looks as if the model m6 performs reasonably as well.

Exercise 3 - now with attitude

Exercise 3, part 1–(VK) 1) Carry on with the model with the two unique intercepts fitted (*scenario* and *subject*) but now build a new model that has *attitude* as a main effect besides *gender*. After create a separate model that besides the main effects of *attitude* and *gender* also include their interaction

```
#making a model with two unicue intercepts and having attitude and gender as main effects
m8 <- lmer(formula = f0mn ~ gender + attitude + (1 | subject) + (1 | scenario), data = pol
iteness)
summary(m8)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: f0mn ~ gender + attitude + (1 | subject) + (1 | scenario)
##    Data: politeness
##
## REML criterion at convergence: 2065.1
##
## Scaled residuals:
##     Min      1Q  Median      3Q     Max
## -2.8511 -0.6081 -0.0602  0.4329  3.8745
##
## Random effects:
##  Groups   Name        Variance Std.Dev.
##  subject  (Intercept) 585.6    24.20
##  scenario (Intercept) 106.7    10.33
##  Residual             882.7    29.71
## Number of obs: 212, groups:  subject, 16; scenario, 7
##
## Fixed effects:
##             Estimate Std. Error t value
## (Intercept)  254.398      9.597  26.507
## genderM     -115.437     12.881  -8.962
## attitudepol  -14.819      4.096  -3.618
##
## Correlation of Fixed Effects:
##             (Intr) gendrM
## genderM     -0.587
## attitudepol -0.220  0.006
```

```
fitted.values(m8)
```

```
##         1         2         3         4         5         6         7         8 
## 216.47929 231.29853 234.27048 249.08973 236.09087 250.91011 241.21606 256.03530 
##         9        10        11        12        13        14        15        16 
## 230.41934 245.23858 219.58594 234.40518 238.54395 223.72471 237.34031 252.15956 
##        17        18        19        20        21        22        23        24 
## 255.13151 269.95075 256.95189 271.77114 262.07708 276.89632 251.28036 266.09960 
##        25        26        27        28        29        30        31        32 
## 240.44696 255.26621 259.40498 244.58573 230.84663 245.66588 248.63783 263.45707 
##        33        34        35        36        37        38        39        40 
## 250.45822 265.27746 255.58340 270.40264 244.78668 259.60593 233.95328 248.77253 
##        41        42        43        44        45        46        47        49 
## 252.91130 238.09206 198.00782 212.82706 215.79901 230.61826 217.61940 222.74458 
##        51        52        53        54        55        57        58        59 
## 211.94787 226.76711 201.11447 215.93371 220.07248 274.39362 289.21286 292.18481 
##        60        61        62        63        64        65        66        67 
## 307.00406 294.00520 308.82444 299.13038 313.94963 288.33367 303.15291 277.50027 
##        68        69        70        71        72        73        74        75 
## 292.31951 296.45828 281.63904 209.70329 224.52253 227.49448 242.31372 229.31487 
##        76        77        79        81        82        83        84        85 
## 244.13411 234.44005 223.64333 212.80994 227.62918 231.76795 216.94871 220.01309 
##        86        87        88        89        90        91        92        93 
## 234.83233 237.80428 252.62353 239.62467 254.44391 244.74985 259.56910 233.95314 
##        94        95        96        97        98        99       100       101 
## 248.77238 223.11974 237.93898 242.07775 227.25851 218.45899 233.27823 236.25019 
##       102       103       104       105       106       107       108       109 
## 251.06943 238.07057 252.88982 243.19576 258.01500 232.39904 247.21828 221.56564 
##       110       111       112       113       114       115       116       117 
## 236.38488 240.52366 225.70441 239.84235 254.66159 257.63354 272.45278 259.45393 
##       118       119       120       121       122       123       124       125 
## 274.27317 264.57911 279.39835 253.78240 268.60164 242.94900 257.76824 261.90701 
##       126       127       128       129       130       131       132       133 
## 247.08777 133.00242 147.82166 150.79361 165.61285 152.61400 167.43324 157.73918 
##       134       135       136       137       138       139       140       141 
## 172.55843 146.94247 161.76171 136.10907 150.92831 155.06708 140.24784 107.79414 
##       142       143       147       149       150       151       152       153 
## 122.61338 125.58533 132.53090 121.73419 136.55343 110.90079 125.72003 129.85880 
##       154       155       156       157       158       159       160       161 
## 115.03956 145.58352 160.40276 163.37471 178.19395 165.19510 180.01434 170.32028 
##       162       163       164       165       166       167       168       169 
## 185.13952 159.52357 174.34281 148.69017 163.50941 167.64818 152.82894 127.46793 
##       170       171       172       173       174       175       176       177 
## 142.28717 145.25913 160.07837 147.07951 161.89876 152.20470 167.02394 141.40798 
##       178       180       181       182       183       184       185       186 
## 156.22722 145.39382 149.53260 134.71335  83.40025  98.21950 101.19145 116.01069 
##       187       188       189       190       191       192       193       194 
## 103.01184 117.83108 108.13702 122.95626  97.34030 112.15955  86.50690 101.32615 
##       195       196       198       200       201       202       203       204 
## 105.46492  90.64568 112.78128 130.57248 117.57362 132.39287 122.69881 137.51805 
##       205       206       207       208       209       210       211       212 
## 111.90209 126.72133 101.06869 115.88793 120.02671 105.20746  87.35321 102.17245 
##       213       214       215       216       217       218       219       220 
## 105.14440 119.96365 106.96479 121.78403 112.08998 126.90922 101.29326 116.11250 
##       221       222       223       224 
##  90.45986 105.27910 109.41787  94.59863
```

```
par(mfrow=c(1,2))
qqnorm(fitted.values(m8))
qqline(fitted.values(m8), col = "red")
qqnorm(politeness$f0mn)
qqline(politeness$f0mn, col = "red")
```

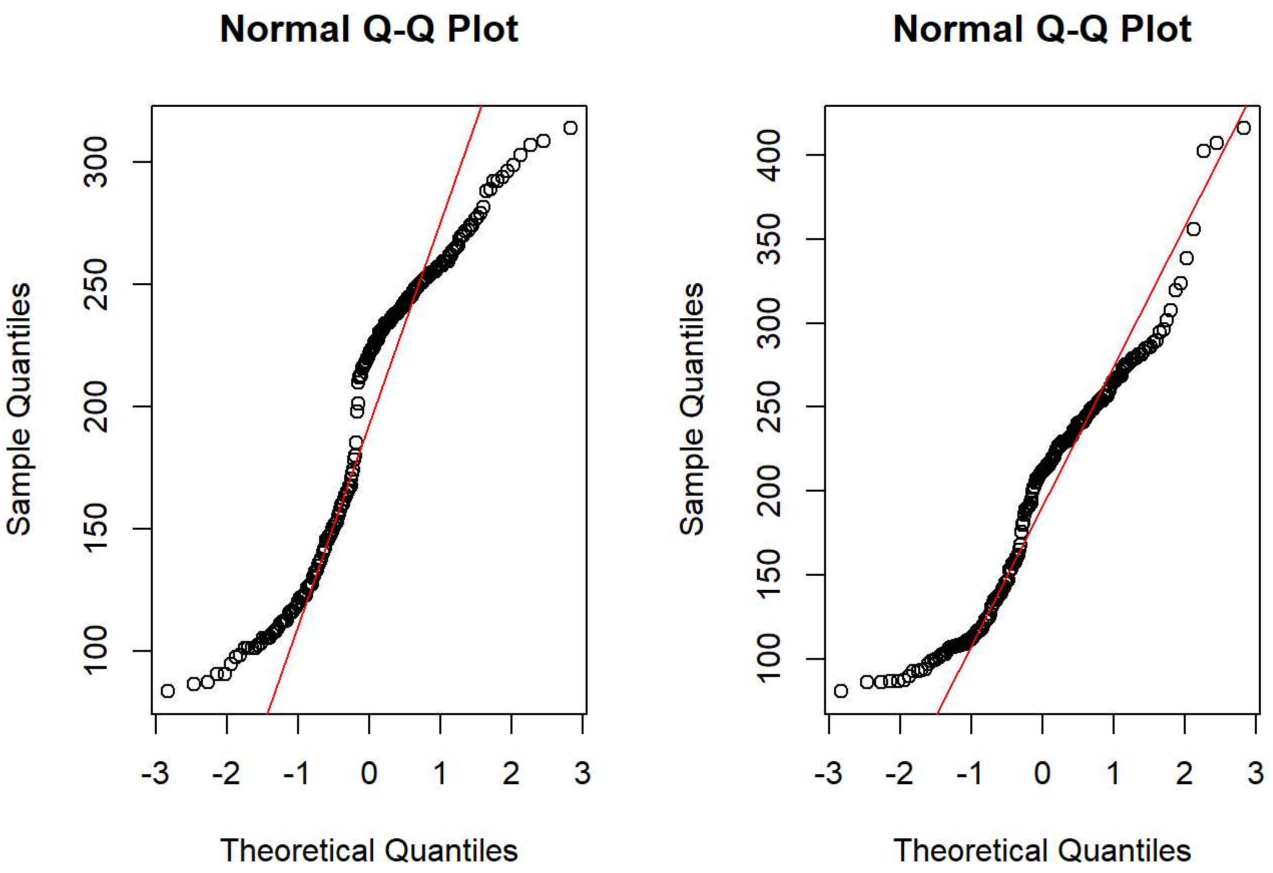


```
plot(m8)
```

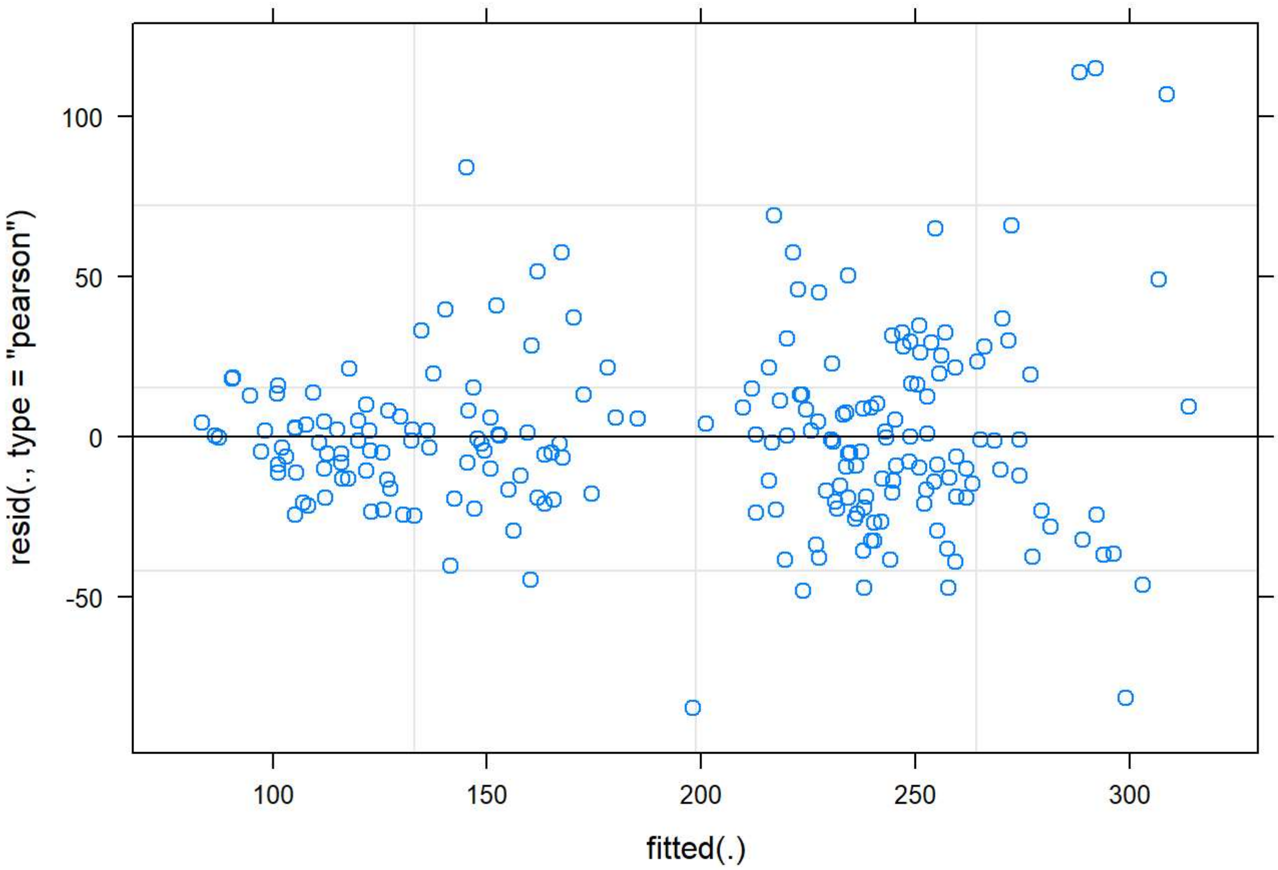


```
#a model that additionally includes the interaction between the main effects (attitude and
gender)
m9 <- lmer(formula = f0mn ~ gender + attitude + gender:attitude + (1 | subject) + (1 | sce
nario), data = politeness)
summary(m9)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: f0mn ~ gender + attitude + gender:attitude + (1 | subject) +
##     (1 | scenario)
##    Data: politeness
##
## REML criterion at convergence: 2058.6
##
## Scaled residuals:
##     Min      1Q  Median      3Q     Max
## -2.8120 -0.5884 -0.0645  0.4014  3.9100
##
## Random effects:
##  Groups   Name        Variance Std.Dev.
##  subject  (Intercept) 584.4    24.17
##  scenario (Intercept) 106.4    10.32
##  Residual             885.5    29.76
## Number of obs: 212, groups:  subject, 16; scenario, 7
##
## Fixed effects:
##                      Estimate Std. Error t value
## (Intercept)           255.618      9.761  26.186
## genderM              -118.232     13.531  -8.738
## attitudepol           -17.192      5.423  -3.170
## genderM:attitudepol     5.544      8.284   0.669
##
## Correlation of Fixed Effects:
##             (Intr) gendrM atttdp
## genderM     -0.606
## attitudepol -0.286  0.206
## gndrM:tttdp  0.187 -0.309 -0.654
```

```
The model m9 can be read as following: The intercept for women/inf are $ \approx 256 Hz$,
when we look at men with the same attitude their pitch drops by $ \approx 118 Hz $. Overal
l a polite attitude will result in a drop in pitch by $ \approx 17 Hz$, however for men it
will only be $ -17.2+5.5 \approx 11.6 Hz$. Korean women's relative drop in pitch is theref
ore larger than male's in a polite situation according to this sample.
```

Exercise 3, part 2–(KV)

2. Compare the three models (1. gender as a main effect; 2. gender and attitude as main effects; 3. gender and attitude as main effects and the interaction between them.

```
#comparing the models
anova(m6, m8, m9)
```

```
## refitting model(s) with ML (instead of REML)
```

```
## Data: politeness
## Models:
## m6: f0mn ~ gender + (1 | subject) + (1 | scenario)
## m8: f0mn ~ gender + attitude + (1 | subject) + (1 | scenario)
## m9: f0mn ~ gender + attitude + gender:attitude + (1 | subject) +
## m9:     (1 | scenario)
##    npar    AIC    BIC  logLik deviance   Chisq Df Pr(>Chisq)
## m6    5 2105.2 2122.0 -1047.6   2095.2
## m8    6 2094.5 2114.6 -1041.2   2082.5 12.6868  1  0.0003683 ***
## m9    7 2096.0 2119.5 -1041.0   2082.0  0.4551  1  0.4998998
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
piecewiseSEM::rsquared(c(m6, m8, m9))
```

```
##   Response   family     link method  Marginal Conditional
## 1     f0mn gaussian identity   none 0.6677206   0.8077964
## 2     f0mn gaussian identity   none 0.6782542   0.8196777
## 3     f0mn gaussian identity   none 0.6782490   0.8192531
```

Exercise 3, part 3–(KV) 3) Choose the model that you think describe the data the best - and write a short report on the main findings based on this model.

```
summary(m8)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: f0mn ~ gender + attitude + (1 | subject) + (1 | scenario)
##    Data: politeness
##
## REML criterion at convergence: 2065.1
##
## Scaled residuals:
##     Min      1Q  Median      3Q     Max
## -2.8511 -0.6081 -0.0602  0.4329  3.8745
##
## Random effects:
##  Groups   Name        Variance Std.Dev.
##  subject  (Intercept) 585.6    24.20
##  scenario (Intercept) 106.7    10.33
##  Residual             882.7    29.71
## Number of obs: 212, groups:  subject, 16; scenario, 7
##
## Fixed effects:
##             Estimate Std. Error t value
## (Intercept)  254.398      9.597  26.507
## genderM     -115.437     12.881  -8.962
## attitudepol  -14.819      4.096  -3.618
##
## Correlation of Fixed Effects:
##             (Intr) gendrM
## genderM     -0.587
## attitudepol -0.220  0.006
```

This dataset consists of the basic demographic information of 16 Korean participants, and their observed pitch in different situations that require either an informal or polite(formal) attitude.

The model we chose as the best is m8. Consistent with the conclusions of the original oauthor's study, this model showed that women on average have a higher pitch than men BUT it also suggested a negative relationship between *attitude* and pitch with p-values<0.001. It would seem that both Korean men's and women's frequency of voice drops when having a polite attitude.

Subjects and scenarios should have different intercepts because it would be assumed they would all have different baselines, (and therefore need different intercepts to account for this). Different subjects will naturally already speak at a different pitch level, so separate intercepts allows us to account for these differences. The different scenarios may also need separate intercepts as certain scenarios may result in participants lowering or raising their pitch to meet the appropriate ambiance of the scenario. Once again, by including separate intercepts for scenarios then we should be accounting for these differences in our models.

Furthermore the output of the summary function shows that more variance is explained by the random effect of the subject than that of the scenario, further strengthening the choice of multilevel modelling.

And here's a QQ-plot of our chosen model

```
qqnorm(fitted.values(m8))
qqline(fitted.values(m8), col = "red")
```